

full_analysis

March 22, 2024

1 Dexian Data Analytics Technical Challenge

Author: James Meredith

Active Project Dates: Mar 15th, 2024 - Mar 22nd, 2024



1.1 Abstract

This technical challenge aims to provide actionable insights from a dataset containing information on American colleges and universities. By applying advanced data analytics techniques, this project seeks to uncover trends and patterns that could assist a consulting firm in advising institutions of higher learning. Through meticulous data analysis, key insights regarding institutional types, acceptance rates, tuition fees, and graduation rates were discovered, enabling data-driven decision-making.

1.2 Business Problem

The scenario for this technical challenge was as follows:

You are working with a firm that provides consulting services for institutions of higher learning.

1.3 The Data

The dataset for this challenge provides detailed information about various American colleges and universities, including public/private designation, application statistics, tuition fees, and graduation rates. The aim was to analyze this dataset to identify trends and insights that could inform strategic advice for higher education institutions.

The dataset used for this technical challenge can be found at the following: <https://docs.google.com/spreadsheets/d/1rThcHm3ZATkhOtsGL6477nQaePrIzcsI/export>

While outside resources could have been considered in order to enrich the initial dataset, due to time constraints and a desire to stay within the bounds of the technical challenge, no outside resources were included in the data pool.

1.3.1 Initial Exploration of The Dataset

Here we ingest the dataset using the Pandas library and perform an initial exploration of the dataset using Pandas, Matplotlib, and Seaborn:

```
[254]: # Imports used in the Analysis
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Ingest the dataset
filepath = './data/Universities.xlsx'
df = pd.read_excel(filepath, sheet_name='usnews3.data.9 .SS (v5.0)')
```

```
[255]: # Generate an overview of the data
df.head()
```

```
[255]:
```

	College Name	State	Public (1)/	Private (2)	\
0	Alaska Pacific University	AK		2	
1	University of Alaska at Fairbanks	AK		1	
2	University of Alaska Southeast	AK		1	
3	University of Alaska at Anchorage	AK		1	
4	Alabama Agri. & Mech. Univ.	AL		1	

	# appli. rec'd	# appl. accepted	# new stud. enrolled	\
0	193.0	146.0	55.0	
1	1852.0	1427.0	928.0	
2	146.0	117.0	89.0	
3	2065.0	1598.0	1162.0	
4	2817.0	1920.0	984.0	

	% new stud. from top 10%	% new stud. from top 25%	# FT undergrad	\
0	16.0	44.0	249.0	

1		NaN		NaN	3885.0
2		4.0		24.0	492.0
3		NaN		NaN	6209.0
4		NaN		NaN	3958.0

	# PT undergrad	in-state tuition	out-of-state tuition	room	board \
0	869.0	7560.0	7560.0	1620.0	2500.0
1	4519.0	1742.0	5226.0	1800.0	1790.0
2	1849.0	1742.0	5226.0	2514.0	2250.0
3	10537.0	1742.0	5226.0	2600.0	2520.0
4	305.0	1700.0	3400.0	1108.0	1442.0

	add. fees	estim. book costs	estim. personal \$	% fac. w/PHD \
0	130.0	800.0	1500.0	76.0
1	155.0	650.0	2304.0	67.0
2	34.0	500.0	1162.0	39.0
3	114.0	580.0	1260.0	48.0
4	155.0	500.0	850.0	53.0

	stud./fac. ratio	Graduation rate
0	11.9	15.0
1	10.0	NaN
2	9.5	39.0
3	13.7	NaN
4	14.3	40.0

```
[256]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1302 entries, 0 to 1301
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   College Name                          1302 non-null   object
1   State                                 1302 non-null   object
2   Public (1)/ Private (2)              1302 non-null   int64
3   # appli. rec'd                       1292 non-null   float64
4   # appl. accepted                     1291 non-null   float64
5   # new stud. enrolled                 1297 non-null   float64
6   % new stud. from top 10%             1067 non-null   float64
7   % new stud. from top 25%             1100 non-null   float64
8   # FT undergrad                       1299 non-null   float64
9   # PT undergrad                       1270 non-null   float64
10  in-state tuition                     1272 non-null   float64
11  out-of-state tuition                 1282 non-null   float64
12  room                                981 non-null    float64
13  board                                804 non-null    float64
14  add. fees                           1028 non-null   float64
```

```

15  estim. book costs      1254 non-null  float64
16  estim. personal $     1121 non-null  float64
17  % fac. w/PHD          1270 non-null  float64
18  stud./fac. ratio       1300 non-null  float64
19  Graduation rate       1204 non-null  float64
dtypes: float64(17), int64(1), object(2)
memory usage: 203.6+ KB

```

```
[257]: df.describe()
```

```

[257]:      Public (1)/ Private (2)  # appli. rec'd  # appl. accepted  \
count      1302.000000      1292.000000      1291.000000
mean        1.639017      2752.097523      1870.683191
std          0.480470      3541.974712      2250.866400
min          1.000000        35.000000        35.000000
25%          1.000000      695.750000      554.500000
50%          2.000000     1470.000000     1095.000000
75%          2.000000     3314.250000     2303.000000
max          2.000000     48094.000000     26330.000000

      # new stud. enrolled  % new stud. from top 10%  \
count      1297.000000      1067.000000
mean        778.880493        25.671978
std        884.578274        18.312618
min         18.000000         1.000000
25%        236.000000        13.000000
50%        447.000000        21.000000
75%        984.000000        32.000000
max       7425.000000        98.000000

      % new stud. from top 25%  # FT undergrad  # PT undergrad  \
count      1100.000000      1299.000000      1270.000000
mean        52.350000      3692.665127      1081.526772
std        20.881316      4544.847897      1672.202912
min         6.000000        59.000000         1.000000
25%        36.750000      966.000000      131.250000
50%        50.000000     1812.000000      472.000000
75%        66.000000     4539.500000     1313.000000
max       100.000000     31643.000000     21836.000000

      in-state tuition  out-of-state tuition      room      board  \
count      1272.000000      1282.000000     981.000000     804.000000
mean      7897.274371      9276.905616     2514.681957     2060.983831
std      5348.162626      4170.770851     1150.836848      661.742099
min       480.000000      1044.000000      500.000000      531.000000
25%      2580.000000      6111.000000     1710.000000     1619.250000
50%      8050.000000      8670.000000     2200.000000     1980.000000

```

75%	11600.000000	11659.000000	3040.000000	2401.500000
max	25750.000000	25750.000000	7400.000000	6250.000000

	add. fees	estim. book costs	estim. personal \$	% fac. w/PHD \
count	1028.000000	1254.000000	1121.000000	1270.000000
mean	392.012646	549.972887	1389.291704	68.645669
std	469.379234	167.355386	714.247857	17.825627
min	9.000000	90.000000	75.000000	8.000000
25%	130.000000	480.000000	900.000000	57.000000
50%	264.500000	502.000000	1250.000000	71.000000
75%	480.000000	600.000000	1794.000000	82.000000
max	4374.000000	2340.000000	6900.000000	105.000000

	stud./fac. ratio	Graduation rate
count	1300.000000	1204.000000
mean	14.858769	60.405316
std	5.186399	18.889058
min	2.300000	8.000000
25%	11.800000	47.000000
50%	14.300000	60.000000
75%	17.600000	74.000000
max	91.800000	118.000000

We note a few things here: - Dataset is 1302 rows (excluding headers) by 20 columns - Dataset contains a mixture of data types (categorical and numeric), some of which are improperly formatted upon ingestion. - Dataset contains null values - Column Headers are a bit difficult to read

And from a cursory look at the summary statistics: - Dataset contains illogical/anomalous values (ie. graduation rates, % of faculty with PhDs both contain values over 100%)

We'll address these issues in the cleaning section of this notebook next.

1.4 Data Cleaning

Here we'll begin the process of cleaning the dataset. Issues addressed are as follows:

- Change Column Names to Improve Readability
- Handle Missing Values
- Handle Duplicate Values
- Remove or Correct Misc. Anomalous Data
- Handle Outliers
- Correct Data Types
- Feature Engineering

[258]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1302 entries, 0 to 1301
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
#   ...
```

```

---  -----
0  College Name          1302 non-null  object
1  State                 1302 non-null  object
2  Public (1)/ Private (2) 1302 non-null  int64
3  # appli. rec'd        1292 non-null  float64
4  # appl. accepted       1291 non-null  float64
5  # new stud. enrolled   1297 non-null  float64
6  % new stud. from top 10% 1067 non-null  float64
7  % new stud. from top 25% 1100 non-null  float64
8  # FT undergrad        1299 non-null  float64
9  # PT undergrad         1270 non-null  float64
10 in-state tuition       1272 non-null  float64
11 out-of-state tuition   1282 non-null  float64
12 room                  981 non-null  float64
13 board                 804 non-null  float64
14 add. fees             1028 non-null  float64
15 estim. book costs     1254 non-null  float64
16 estim. personal $     1121 non-null  float64
17 % fac. w/PHD          1270 non-null  float64
18 stud./fac. ratio       1300 non-null  float64
19 Graduation rate       1204 non-null  float64
dtypes: float64(17), int64(1), object(2)
memory usage: 203.6+ KB

```

1.4.1 Changing Column Names To Improve Readability

Firstly, Column Names have been changed for improved readability

```

[259]: # Renaming columns for improved readability
df.rename(columns={
    'College Name': 'college_name',
    'State': 'state',
    'Public (1)/ Private (2)': 'institution_type',
    '# appli. rec\'d': 'applications_received',
    '# appl. accepted': 'applications_accepted',
    '# new stud. enrolled': 'new_students_enrolled',
    '% new stud. from top 10%': 'percent_from_top_10',
    '% new stud. from top 25%': 'percent_from_top_25',
    '# FT undergrad': 'full_time_undergrads',
    '# PT undergrad': 'part_time_undergrads',
    'in-state tuition': 'in_state_tuition',
    'out-of-state tuition': 'out_of_state_tuition',
    'room': 'room_costs',
    'board': 'board_costs',
    'add. fees': 'additional_fees',
    'estim. book costs': 'estimated_book_costs',
    'estim. personal $': 'estimated_personal_expenses',
    '% fac. w/PHD': 'percent_faculty_with_phd',

```

```

    'stud./fac. ratio': 'student_faculty_ratio',
    'Graduation rate': 'graduation_rate',
}, inplace=True)

```

```
[260]: df.columns
```

```

[260]: Index(['college_name', 'state', 'institution_type', 'applications_received',
            'applications_accepted', 'new_students_enrolled', 'percent_from_top_10',
            'percent_from_top_25', 'full_time_undergrads', 'part_time_undergrads',
            'in_state_tuition', 'out_of_state_tuition', 'room_costs', 'board_costs',
            'additional_fees', 'estimated_book_costs',
            'estimated_personal_expenses', 'percent_faculty_with_phd',
            'student_faculty_ratio', 'graduation_rate'],
            dtype='object')

```

1.4.2 Handling Null Values

Next, here we discuss various strategies for handling missing values in the dataset.

```

[261]: # Calculate percentage of missing values per column
missing_percentage = (df.isnull().sum() / len(df)) * 100
print(missing_percentage)

```

```

college_name      0.000000
state             0.000000
institution_type   0.000000
applications_received  0.768049
applications_accepted  0.844854
new_students_enrolled  0.384025
percent_from_top_10  18.049155
percent_from_top_25  15.514593
full_time_undergrads  0.230415
part_time_undergrads  2.457757
in_state_tuition    2.304147
out_of_state_tuition  1.536098
room_costs         24.654378
board_costs        38.248848
additional_fees     21.044547
estimated_book_costs  3.686636
estimated_personal_expenses 13.901690
percent_faculty_with_phd  2.457757
student_faculty_ratio  0.153610
graduation_rate     7.526882
dtype: float64

```

Possible strategies for handling Null Values in the dataset are as follows: - Removing Null values from the dataset - Filling in Null values (by imputing the mean, median, using KNNImputer, etc.) - Setting Nulls = 0 - De-emphasis or Elimination of Column from Analysis Entirely

Next we'll discuss the pros and cons of each method, and discuss how we'll apply them to the dataset.

1. Removing Null values from the dataset - Pros: - Ensures Data Purity and avoids the introduction of bias - Simplicity - Cons: - Data Loss and Reduction in Statistical Power: Continued removal of data from the dataset can limit the number of overall observations and reduce statistical power

2. Filling in Null Values: - Pros: - Data Preservation: We maximize the preservation of the overall dataset - Cons: - Possible Introduction of Bias: If the imputation method does not align well with the true data distribution or the reason behind the missingness, we risk introducing bias into the dataset. - Reduced variability: Replacing missing values with the median can artificially reduce the variability in the dataset, leading to an underestimate in the observed variance, covariance, etc.

3. Setting Nulls = 0: - Depending on the method of data collection, null values here might be indicating a true value of zero. Because it was unclear from this dataset, the author chose not to use this method when handling null values.

4. De-emphasis or Elimination of Column from Analysis Entirely: - For columns with a high percentage of missing values, it may be wise to eliminate them from the analysis entirely, or simply leave the column as-is and acknowledge a lack of data during the final analysis.

To decide how to apply each strategy to the dataset, we'll categorize each column according to its percentage of missing values and assign one of the above strategies accordingly. Categories are as follows: - Columns with 0% Missing Values - Columns with a Small Percentage of Missing Values (<5%) - Columns with a Moderate Percentage of Missing Values (5-20%) - Columns with a High Amount of Missing Values (>20%)

Columns with 0% Missing Values - unchanged

```
college_name
state
institution_type
percent_faculty_with_phd
student_faculty_ratio
graduation_rate
```

Columns with a Small Percentage of Missing Values (<5%) - We'll remove these from the dataset to maximize data purity, although filling in null values via imputation would have also been appropriate here.

```
applications_received - 0.511073%
applications_accepted - 0.425894%
new_students_enrolled - 0.170358%
full_time_undergrads - 0.255537%
part_time_undergrads - 2.470187%
in_state_tuition - 2.129472%
out_of_state_tuition - 1.277683%
estimated_book_costs - 2.896082%
percent_faculty_with_phd - 2.457757%
```

Columns with a Moderate Percentage of Missing Values (5-20%) - For percent_from_top_10 and percent_from_top_25, we'll make the assumption that missing values should be properly inter-

puted as 0, meaning no students from the top of high school classes attended these universities, although confirmation of this from an enhanced understanding of the data gathering process would be desirable. For `estimated_personal_expenses` and `additional_fees`, we'll impute the median (which is more resistant to skewness in the distribution in the data than the mean), as it's assumed that every university will have some additional fees associated with a program, and every student will face some out-of-pocket costs to attend a university, even if only minimally. For graduation rate, it's assumed that every accredited university has at level some level of graduation, and so the median will be imputed for missing values rather than setting them to 0.

```
percent_from_top_10 - 15.587734%
percent_from_top_25 - 13.117547%
estimated_personal_expenses - 12.095400%
additional_fees - 20.528109%
graduation_rate - 7.526882%
```

Columns with a High Amount of Missing Values (>20%) - Due to the high number of missing values and data integrity concerns, we'll leave these columns as-is be de-emphasis them from the analysis.

```
room_costs - 23.594549%
board_costs - 37.223169%
```

```
[262]: # Drop Missing values from columns with small percentage of null values
df.dropna(subset=['applications_received', 'applications_accepted',
↳ 'new_students_enrolled', 'full_time_undergrads', 'part_time_undergrads',
↳ 'in_state_tuition', 'out_of_state_tuition', 'estimated_book_costs', 'percent_faculty_with_phd',
↳ inplace=True])
```

```
[263]: # Fill percent_from_top_10 and percent_from_top_25 with 0
df['percent_from_top_10'].fillna(0, inplace=True)
df['percent_from_top_25'].fillna(0, inplace=True)
```

```
[264]: # Fill estimated_personal_expenses and additional_fees with median
df['estimated_personal_expenses'].fillna(df['estimated_personal_expenses'].
↳ median(), inplace=True)
df['additional_fees'].fillna(df['additional_fees'].median(), inplace=True)
df['graduation_rate'].fillna(df['graduation_rate'].median(), inplace=True)
```

```
[265]: # Calculate percentage of missing values per column
missing_percentage = (df.isnull().sum() / len(df)) * 100
print(missing_percentage)
```

```
college_name          0.000000
state                 0.000000
institution_type      0.000000
applications_received 0.000000
applications_accepted 0.000000
new_students_enrolled 0.000000
percent_from_top_10   0.000000
percent_from_top_25   0.000000
```

```

full_time_undergrads      0.000000
part_time_undergrads      0.000000
in_state_tuition          0.000000
out_of_state_tuition      0.000000
room_costs                22.723404
board_costs                37.191489
additional_fees            0.000000
estimated_book_costs      0.000000
estimated_personal_expenses 0.000000
percent_faculty_with_phd  0.000000
student_faculty_ratio     0.000000
graduation_rate           0.000000
dtype: float64

```

1.4.3 Handling Duplicate Data

Next we check for and handle duplicates in the dataset.

```
[266]: duplicate_rows = df.duplicated()
print(f"Number of duplicate rows: {duplicate_rows.sum()}")
```

Number of duplicate rows: 0

After cleaning missing values the dataset appears to have no duplicate values, and so we move on to the next stage of cleaning the dataset.

1.4.4 Removing Misc. Anomalous Data

Upon our initial inspection of the dataset we noted some anomalous data (eg. `graduation_rate` and `percent_faculty_with_phd` both contain values over 100%). Anomalous data is expunged from the dataset here. We begin by examining summary statistics to check for illogical or impossible values like those noted above, and then perform some cross-field validation to check for logical inconsistencies (ex. in-state tuitions for a university that are higher than out of state, etc.)

```
[267]: df.describe()
```

```
[267]:
```

	institution_type	applications_received	applications_accepted	\
count	1175.000000	1175.000000	1175.000000	
mean	1.651064	2663.513191	1839.102979	
std	0.476837	3408.747132	2187.012227	
min	1.000000	52.000000	36.000000	
25%	1.000000	684.000000	553.000000	
50%	2.000000	1450.000000	1086.000000	
75%	2.000000	3286.500000	2309.000000	
max	2.000000	48094.000000	26330.000000	

	new_students_enrolled	percent_from_top_10	percent_from_top_25	\
count	1175.000000	1175.000000	1175.000000	
mean	769.183830	20.908936	44.501277	

std	881.649886	18.098567	26.108308
min	18.000000	0.000000	0.000000
25%	234.000000	9.000000	29.000000
50%	444.000000	18.000000	46.000000
75%	945.000000	29.000000	63.000000
max	7425.000000	96.000000	100.000000

	full_time_undergrads	part_time_undergrads	in_state_tuition \
count	1175.000000	1175.000000	1175.000000
mean	3642.939574	1047.977021	7906.854468
std	4491.340232	1649.297896	5273.593853
min	88.000000	1.000000	480.000000
25%	967.500000	128.000000	2615.500000
50%	1803.000000	442.000000	8190.000000
75%	4475.000000	1283.000000	11515.000000
max	31643.000000	21836.000000	25750.000000

	out_of_state_tuition	room_costs	board_costs	additional_fees \
count	1175.000000	908.000000	738.000000	1175.000000
mean	9297.880851	2515.396476	2044.966125	339.243404
std	4094.042844	1151.534021	637.178853	301.456287
min	1044.000000	500.000000	531.000000	9.000000
25%	6161.000000	1710.750000	1620.500000	150.000000
50%	8734.000000	2200.000000	1980.000000	260.000000
75%	11652.000000	3000.000000	2392.250000	400.000000
max	25750.000000	7400.000000	6250.000000	3247.000000

	estimated_book_costs	estimated_personal_expenses \
count	1175.000000	1175.000000
mean	545.386383	1370.233191
std	159.317616	680.727562
min	90.000000	250.000000
25%	475.500000	920.000000
50%	500.000000	1230.000000
75%	600.000000	1700.000000
max	2340.000000	6900.000000

	percent_faculty_with_phd	student_faculty_ratio	graduation_rate
count	1175.000000	1175.000000	1175.000000
mean	68.740426	14.879830	60.549787
std	17.617501	5.245299	17.965015
min	8.000000	2.300000	8.000000
25%	57.000000	11.800000	49.000000
50%	71.000000	14.300000	60.000000
75%	82.000000	17.500000	73.000000
max	103.000000	91.800000	118.000000

We note two instances of anomalies - percent_faculty_with_phd greater than 100%, and grad-

uation_rate over 100%. Because we cannot be certain whether these values were intended to be 100% or some other value entirely, we remove them from the dataset here.

```
[268]: # Example: Removing rows where 'Graduation rate' is above 100%
```

```
df = df[df['percent_faculty_with_phd'] <= 100]
df = df[df['graduation_rate'] <= 100]
```

```
[269]: df.describe()
```

```
[269]:
```

	institution_type	applications_received	applications_accepted	\
count	1173.000000	1173.000000	1173.000000	
mean	1.651321	2664.323956	1838.901961	
std	0.476755	3410.909317	2188.022662	
min	1.000000	52.000000	36.000000	
25%	1.000000	686.000000	553.000000	
50%	2.000000	1450.000000	1086.000000	
75%	2.000000	3281.000000	2306.000000	
max	2.000000	48094.000000	26330.000000	

	new_students_enrolled	percent_from_top_10	percent_from_top_25	\
count	1173.000000	1173.000000	1173.000000	
mean	769.838875	20.918159	44.507246	
std	882.239351	18.110632	26.128999	
min	18.000000	0.000000	0.000000	
25%	233.000000	9.000000	29.000000	
50%	444.000000	18.000000	46.000000	
75%	946.000000	29.000000	63.000000	
max	7425.000000	96.000000	100.000000	

	full_time_undergrads	part_time_undergrads	in_state_tuition	\
count	1173.000000	1173.000000	1173.000000	
mean	3647.261722	1049.639386	7911.670929	
std	4493.947007	1650.210361	5273.805830	
min	88.000000	1.000000	480.000000	
25%	967.000000	128.000000	2625.000000	
50%	1803.000000	446.000000	8190.000000	
75%	4481.000000	1283.000000	11520.000000	
max	31643.000000	21836.000000	25750.000000	

	out_of_state_tuition	room_costs	board_costs	additional_fees	\
count	1173.000000	906.000000	736.000000	1173.000000	
mean	9301.590793	2516.556291	2045.112772	339.212276	
std	4095.481045	1152.363247	637.646380	301.708472	
min	1044.000000	500.000000	531.000000	9.000000	
25%	6172.000000	1711.250000	1621.500000	150.000000	
50%	8734.000000	2200.000000	1980.000000	260.000000	
75%	11656.000000	3000.000000	2390.750000	400.000000	

max	25750.000000	7400.000000	6250.000000	3247.000000
-----	--------------	-------------	-------------	-------------

	estimated_book_costs	estimated_personal_expenses \
count	1173.000000	1173.000000
mean	545.293265	1371.589088
std	159.437506	680.507288
min	90.000000	250.000000
25%	475.000000	930.000000
50%	500.000000	1230.000000
75%	600.000000	1700.000000
max	2340.000000	6900.000000

	percent_faculty_with_phd	student_faculty_ratio	graduation_rate
count	1173.000000	1173.000000	1173.000000
mean	68.751066	14.878176	60.515772
std	17.551079	5.249229	17.894480
min	8.000000	2.300000	8.000000
25%	57.000000	11.800000	49.000000
50%	71.000000	14.200000	60.000000
75%	82.000000	17.500000	73.000000
max	100.000000	91.800000	100.000000

```
[270]: # Check if in-state tuition is ever higher than out-of-state tuition
inconsistent_tuition = df[df['in_state_tuition'] > df['out_of_state_tuition']]
print(f"Number of records where in-state tuition is higher than out-of-state_
      ↪tuition: {len(inconsistent_tuition)}")

# Keep only the records where in-state tuition is less than or equal to_
      ↪out-of-state tuition
df = df[df['in_state_tuition'] <= df['out_of_state_tuition']]

# Checking if accepted applications exceed received applications
inconsistent_applications = df[df['applications_received'] <_
      ↪df['applications_accepted']]
print(f"Number of records where accepted applications exceed received_
      ↪applications: {len(inconsistent_applications)}")
```

Number of records where in-state tuition is higher than out-of-state tuition: 2
 Number of records where accepted applications exceed received applications: 0

1.4.5 Handling Outliers

In this section we identify and discuss strategies for dealing with outliers.

Possible Strategies for Dealing with Outliers - Removal - Capping - Transformation - Leaving as Is

```
[271]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1171 entries, 0 to 1301
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   college_name                          1171 non-null   object
1   state                                 1171 non-null   object
2   institution_type                       1171 non-null   int64
3   applications_received                 1171 non-null   float64
4   applications_accepted                 1171 non-null   float64
5   new_students_enrolled                 1171 non-null   float64
6   percent_from_top_10                  1171 non-null   float64
7   percent_from_top_25                  1171 non-null   float64
8   full_time_undergrads                 1171 non-null   float64
9   part_time_undergrads                 1171 non-null   float64
10  in_state_tuition                      1171 non-null   float64
11  out_of_state_tuition                  1171 non-null   float64
12  room_costs                           904 non-null    float64
13  board_costs                           734 non-null    float64
14  additional_fees                       1171 non-null   float64
15  estimated_book_costs                  1171 non-null   float64
16  estimated_personal_expenses           1171 non-null   float64
17  percent_faculty_with_phd              1171 non-null   float64
18  student_faculty_ratio                 1171 non-null   float64
19  graduation_rate                       1171 non-null   float64
dtypes: float64(17), int64(1), object(2)
memory usage: 192.1+ KB

```

```

[272]: # Create Visualizations to check for Outliers

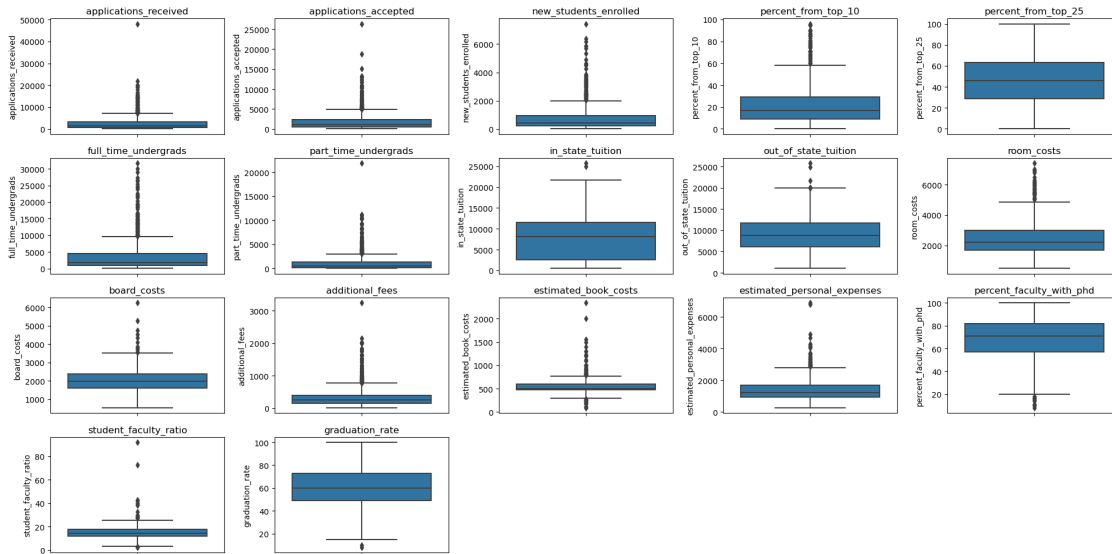
# Selecting numeric columns to visualize. Exclude 'institution_type' since it's
↳ categorical represented as int64.
numeric_cols = df.select_dtypes(include=['float64']).columns

# Create a large figure to accommodate the subplots
plt.figure(figsize=(20, 10))

# Create a boxplot for each numeric column
for index, column in enumerate(numeric_cols, 1):
    plt.subplot(4, 5, index) # Adjust grid dimensions (4x5) based on the
↳ number of columns
    sns.boxplot(y=df[column])
    plt.title(column)
    plt.tight_layout()

plt.show()

```



A number of columns appear to have numerous outliers, we attempt to get a value count for each column below:

```
[273]: outlier_counts = {}

# Loop through each column in the DataFrame
for column in df.select_dtypes(include=['float64', 'int64']).columns:
    # Calculate Q1 and Q3, and then IQR
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    # Define outliers as values below (Q1 - 1.5 * IQR) or above (Q3 + 1.5 * IQR)
    outliers = df[(df[column] < (Q1 - 1.5 * IQR)) | (df[column] > (Q3 + 1.5 * IQR))]
    # Count the number of outliers
    outlier_counts[column] = outliers.shape[0]

# Print the number of outliers in each column
for column, count in outlier_counts.items():
    print(f"{column}: {count} outliers")
```

```
institution_type: 0 outliers
applications_received: 102 outliers
applications_accepted: 92 outliers
new_students_enrolled: 94 outliers
percent_from_top_10: 48 outliers
percent_from_top_25: 0 outliers
full_time_undergrads: 114 outliers
part_time_undergrads: 105 outliers
```

```

in_state_tuition: 2 outliers
out_of_state_tuition: 6 outliers
room_costs: 44 outliers
board_costs: 11 outliers
additional_fees: 93 outliers
estimated_book_costs: 81 outliers
estimated_personal_expenses: 41 outliers
percent_faculty_with_phd: 10 outliers
student_faculty_ratio: 21 outliers
graduation_rate: 2 outliers

```

While handling of outliers should ordinarily be done with great care and attention paid to the goals of the analysis and their affect on the data’s distribution, due to the time constraints of this analysis we will retain them in the dataset and discuss them as they come up in the analysis section of this notebook.

1.4.6 Correcting Data Types

Here we improve the formatting of several columns. For increased readability, we’ll first convert the institution type (Public/Private) from a numeric representation to “Public” or “Private”. Additionally, several columns referencing number of students, faculty, and applications, where decimal values don’t make sense, are changed to integer data type here as well.

[274]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1171 entries, 0 to 1301
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   college_name          1171 non-null   object
 1   state                 1171 non-null   object
 2   institution_type       1171 non-null   int64
 3   applications_received  1171 non-null   float64
 4   applications_accepted  1171 non-null   float64
 5   new_students_enrolled  1171 non-null   float64
 6   percent_from_top_10    1171 non-null   float64
 7   percent_from_top_25    1171 non-null   float64
 8   full_time_undergrads   1171 non-null   float64
 9   part_time_undergrads   1171 non-null   float64
10   in_state_tuition       1171 non-null   float64
11   out_of_state_tuition   1171 non-null   float64
12   room_costs            904 non-null    float64
13   board_costs           734 non-null    float64
14   additional_fees        1171 non-null   float64
15   estimated_book_costs   1171 non-null   float64
16   estimated_personal_expenses  1171 non-null   float64
17   percent_faculty_with_phd  1171 non-null   float64
18   student_faculty_ratio  1171 non-null   float64

```



```

19  graduation_rate          1171 non-null   float64
dtypes: float64(17), int64(1), object(2)
memory usage: 192.1+ KB

```

```

[275]: # Define a mapping dictionary to translate 1 to 'Public' and 2 to 'Private'
institution_type_map = {1: 'Public', 2: 'Private'}

# Use the map function to apply this translation to the 'Public (1)/ Private_
↪(2)' column
df['institution_type'] = df['institution_type'].map(institution_type_map)

```

```

[276]: # Correcting Data Types
data_types = {
    'institution_type' : 'object',
    'applications_received' : 'int64',
    'applications_accepted' : 'int64',
    'new_students_enrolled' : 'int64',
    'full_time_undergrads' : 'int64',
    'part_time_undergrads' : 'int64',
}

# Apply the data type changes to your DataFrame
df = df.astype(data_types)

```

```

[277]: df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1171 entries, 0 to 1301
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   college_name          1171 non-null   object
 1   state                 1171 non-null   object
 2   institution_type       1171 non-null   object
 3   applications_received  1171 non-null   int64
 4   applications_accepted  1171 non-null   int64
 5   new_students_enrolled 1171 non-null   int64
 6   percent_from_top_10    1171 non-null   float64
 7   percent_from_top_25    1171 non-null   float64
 8   full_time_undergrads   1171 non-null   int64
 9   part_time_undergrads   1171 non-null   int64
10   in_state_tuition       1171 non-null   float64
11   out_of_state_tuition   1171 non-null   float64
12   room_costs            904 non-null    float64
13   board_costs           734 non-null    float64
14   additional_fees        1171 non-null   float64
15   estimated_book_costs   1171 non-null   float64
16   estimated_personal_expenses 1171 non-null   float64

```

```

17 percent_faculty_with_phd      1171 non-null    float64
18 student_faculty_ratio         1171 non-null    float64
19 graduation_rate               1171 non-null    float64
dtypes: float64(12), int64(5), object(3)
memory usage: 192.1+ KB

```

1.4.7 Feature Engineering

Here we create a few columns containing metrics helpful to our analysis, namely Acceptance Rate, Yield Rate, Total Cost of Attendance, Total Enrollment

```

[278]: # Create Aforementioned Columns
df['acceptance_rate'] = (df['applications_accepted'] /
    ↪df['applications_received']) * 100
df['yield_rate'] = (df['new_students_enrolled'] / df['applications_accepted'])
    ↪* 100
df['total_cost_in_state'] = df['in_state_tuition'] + df['room_costs'].fillna(0)
    ↪+ df['board_costs'].fillna(0) + df['additional_fees'] +
    ↪df['estimated_book_costs'] + df['estimated_personal_expenses']
df['total_enrollment'] = df['full_time_undergrads'] + df['part_time_undergrads']

```

```

[279]: df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1171 entries, 0 to 1301
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   college_name                          1171 non-null   object
1   state                                 1171 non-null   object
2   institution_type                      1171 non-null   object
3   applications_received                 1171 non-null   int64
4   applications_accepted                 1171 non-null   int64
5   new_students_enrolled                 1171 non-null   int64
6   percent_from_top_10                  1171 non-null   float64
7   percent_from_top_25                  1171 non-null   float64
8   full_time_undergrads                 1171 non-null   int64
9   part_time_undergrads                 1171 non-null   int64
10  in_state_tuition                      1171 non-null   float64
11  out_of_state_tuition                  1171 non-null   float64
12  room_costs                           904 non-null    float64
13  board_costs                           734 non-null    float64
14  additional_fees                       1171 non-null   float64
15  estimated_book_costs                  1171 non-null   float64
16  estimated_personal_expenses           1171 non-null   float64
17  percent_faculty_with_phd             1171 non-null   float64
18  student_faculty_ratio                 1171 non-null   float64
19  graduation_rate                      1171 non-null   float64

```

```

20 acceptance_rate          1171 non-null    float64
21 yield_rate                1171 non-null    float64
22 total_cost_in_state       1171 non-null    float64
23 total_enrollment          1171 non-null    int64
dtypes: float64(15), int64(6), object(3)
memory usage: 228.7+ KB

```

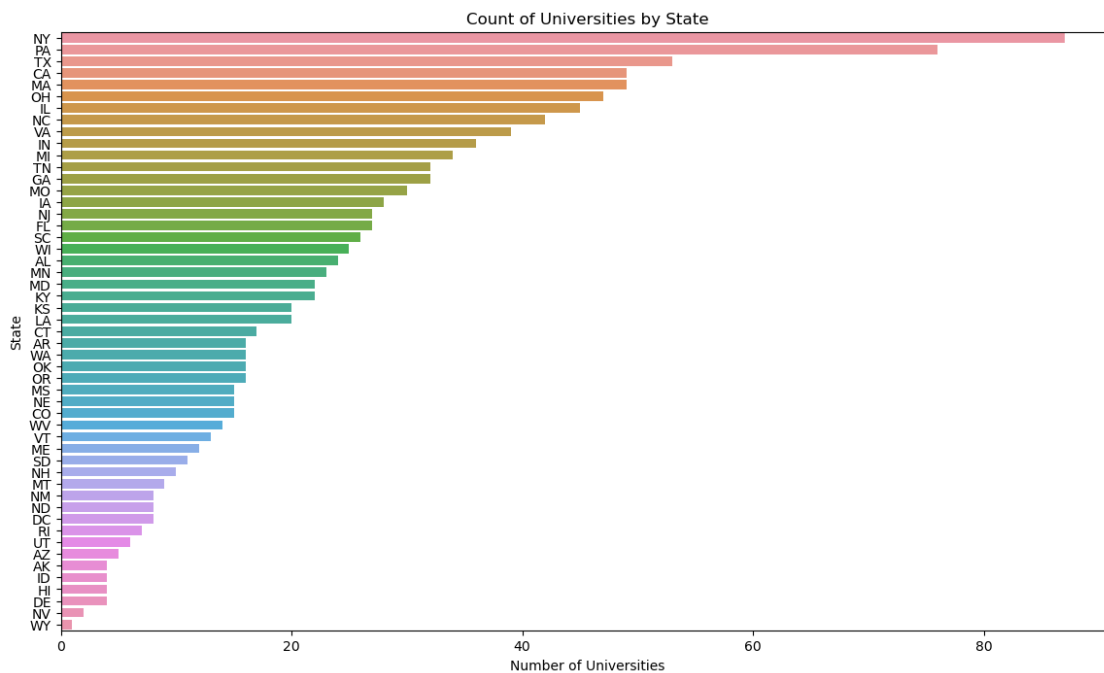
1.5 Exploratory Data Analysis/Results

In this section, we analyze our cleaned prepared dataset

```

[280]: # Visualization: Distribution of Universities by State
plt.figure(figsize=(14, 8))
sns.countplot(y='state', data=df, order = df['state'].value_counts().index)
plt.title('Count of Universities by State')
plt.xlabel('Number of Universities')
plt.ylabel('State')
plt.show()

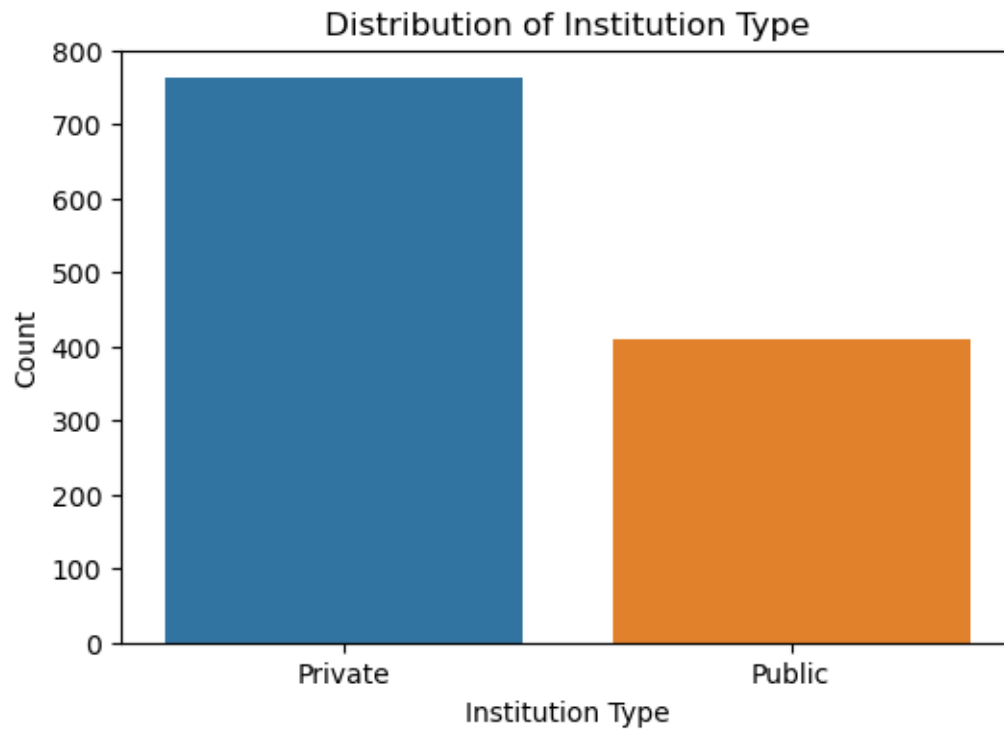
```



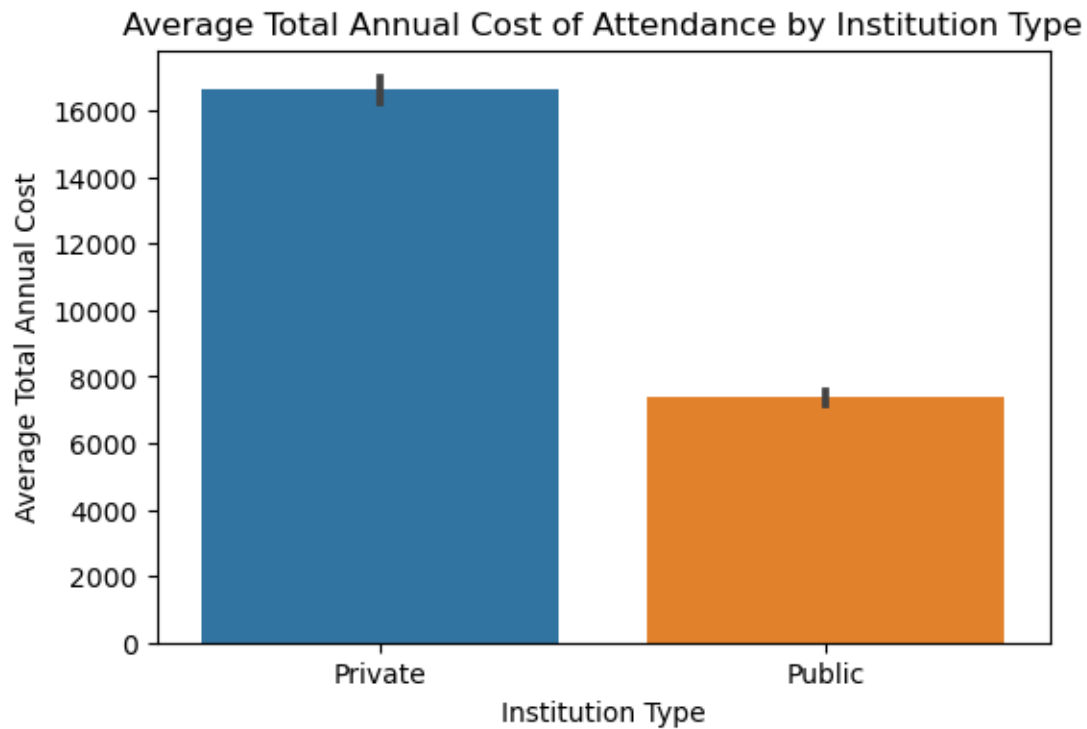
```

[281]: # Visualization: Distribution of public vs. private institutions
plt.figure(figsize=(6, 4))
sns.countplot(x='institution_type', data=df)
plt.title('Distribution of Institution Type')
plt.xlabel('Institution Type')
plt.ylabel('Count')
plt.show()

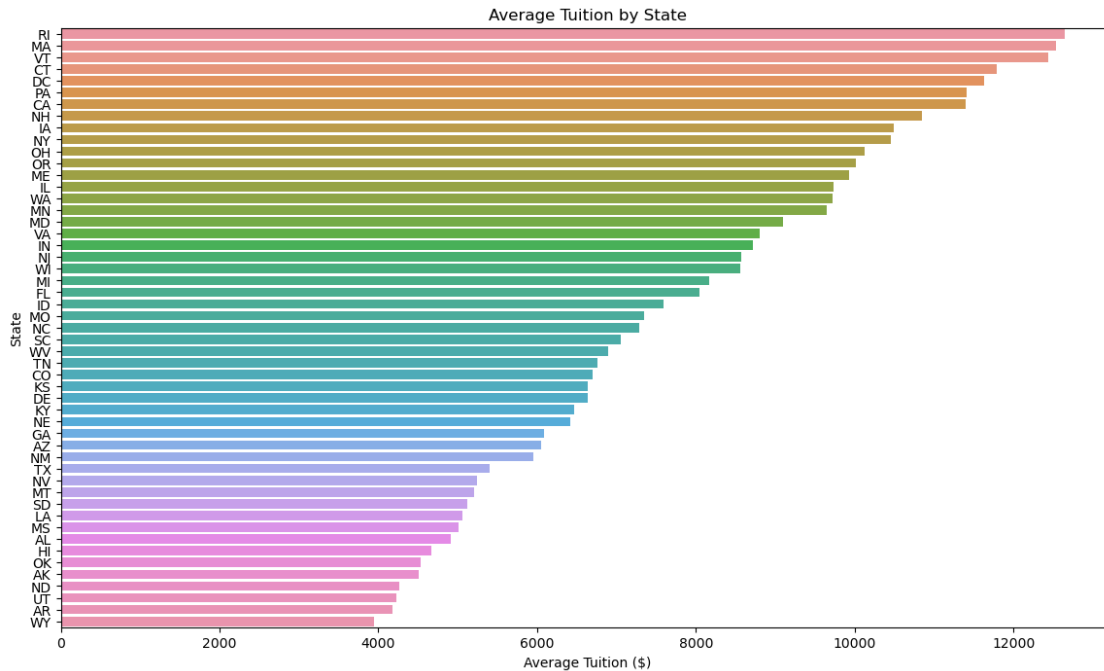
```




```
[282]: # Visualization: Average Total Annual Cost of Attendance for Public vs Private
↳ Institutions
df['total_cost'] = df['total_cost_in_state']
plt.figure(figsize=(6, 4))
sns.barplot(x='institution_type', y='total_cost', data=df, estimator=np.mean)
plt.title('Average Total Annual Cost of Attendance by Institution Type')
plt.xlabel('Institution Type')
plt.ylabel('Average Total Annual Cost')
plt.show()
```



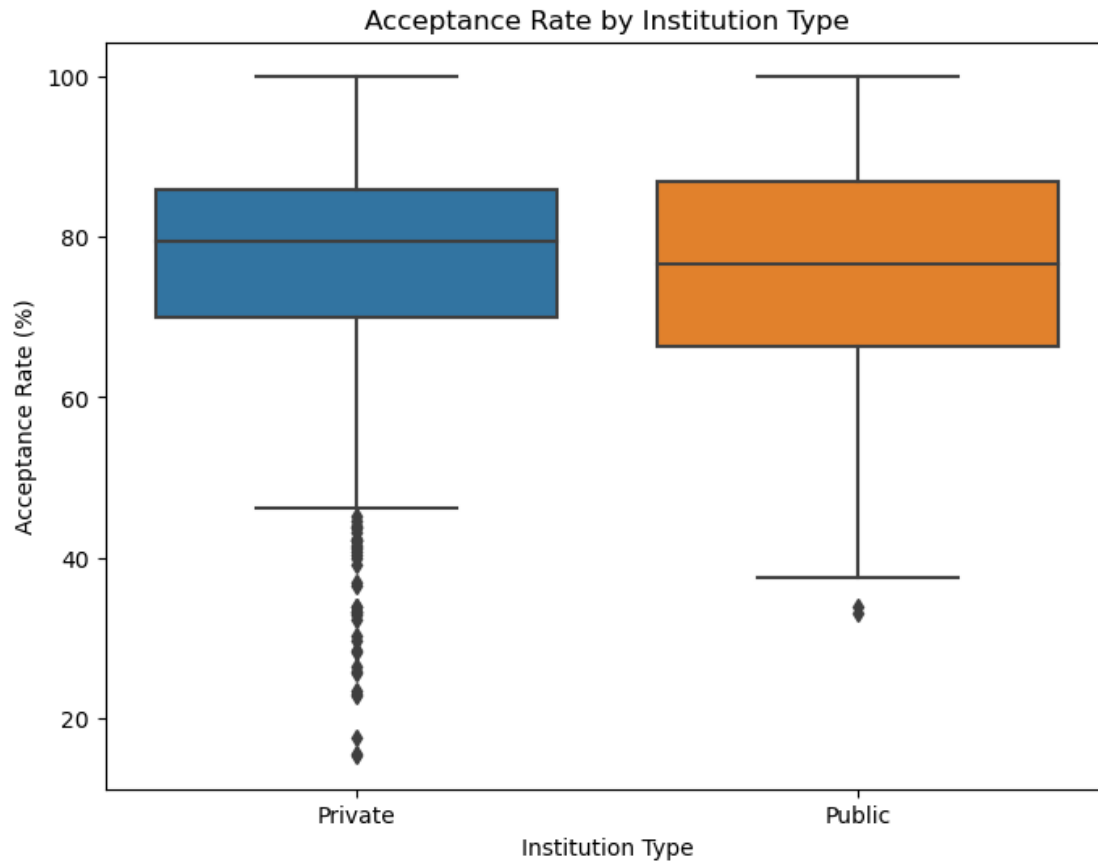
```
[ ]: # Visualization: Average Total Tuition Cost by State
plt.figure(figsize=(14, 8))
df['average_tuition'] = (df['in-state tuition'] + df['out-of-state tuition']) / 2
state_tuition = df.groupby('State')['average_tuition'].mean().
    sort_values(ascending=False)
sns.barplot(x=state_tuition.values, y=state_tuition.index)
plt.title('Average Tuition by State')
plt.xlabel('Average Tuition ($)')
plt.ylabel('State')
plt.show()
```



```
[287]: # Visualization: Distribution of Tuition Fees
plt.figure(figsize=(12, 6))
sns.histplot(df['in_state_tuition'], kde=True, color='skyblue', label='In-State Tuition')
sns.histplot(df['out_of_state_tuition'], kde=True, color='red', label='Out-of-State Tuition')
plt.title('Distribution of Tuition Fees')
plt.xlabel('Tuition Fee ($)')
plt.ylabel('Count')
plt.legend()
plt.show()
```




```
[291]: # Visualization: Distribution of Acceptance Rates by Institution Type
plt.figure(figsize=(8, 6))
sns.boxplot(x='institution_type', y='acceptance_rate', data=df)
plt.title('Acceptance Rate by Institution Type')
plt.xlabel('Institution Type')
plt.ylabel('Acceptance Rate (%)')
plt.show()
```




```
[283]: # Visualization: Average Student Body Composition (Full-time vs. Part-time)
df['total_students'] = df['full_time_undergrads'] + df['part_time_undergrads']
df['percent_full_time'] = (df['full_time_undergrads'] / df['total_students']) * 100

plt.figure(figsize=(6, 4))
sns.barplot(x='institution_type', y='percent_full_time', data=df, estimator=np.mean)

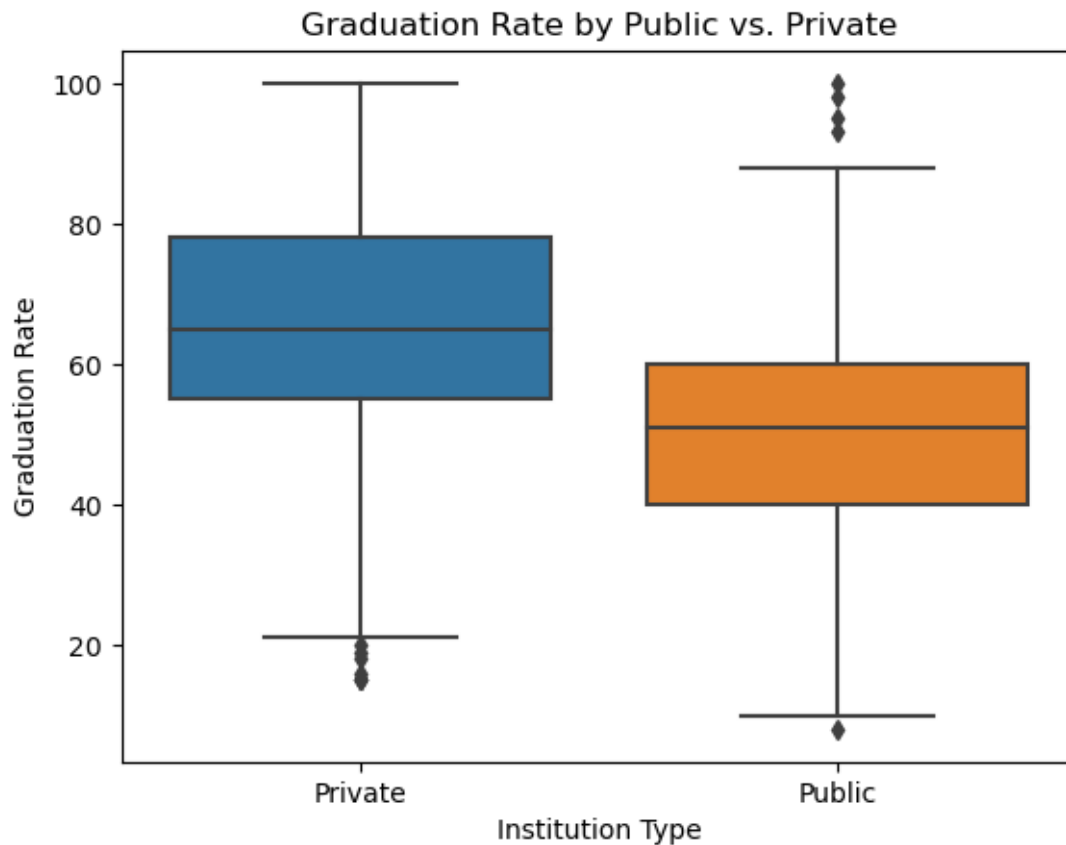
plt.title('Average Percentage of Full-Time Students by Institution Type')
plt.xlabel('Institution Type')
plt.ylabel('Percentage of Full-Time Students')
plt.show()
```

```
[288]: # Visualization: Distribution of Graduation Rates
plt.figure(figsize=(10, 6))
sns.histplot(df['graduation_rate'], bins=20, kde=True)
plt.title('Distribution of Graduation Rates')
plt.xlabel('Graduation Rate (%)')
plt.ylabel('Frequency')
plt.show()
```

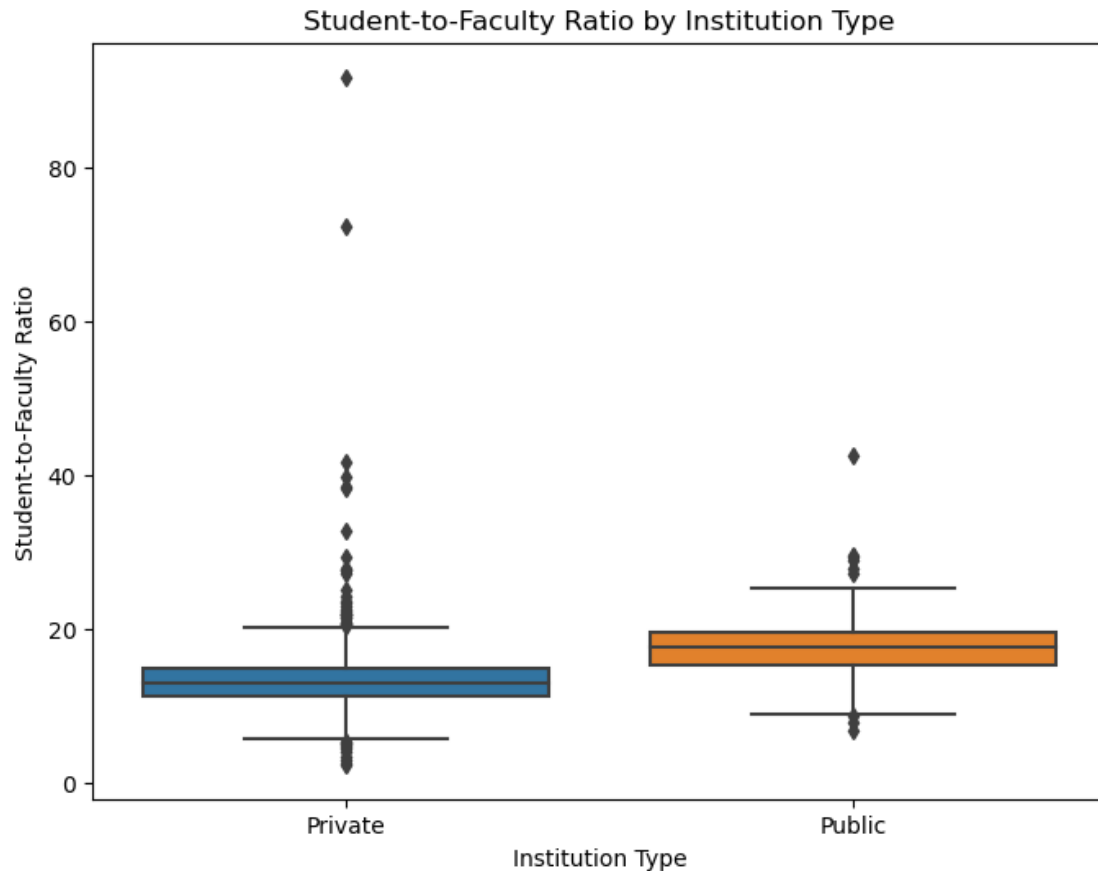


```
[284]: # Visualization: Average Student Body Composition (Full-time vs. Part-time)
sns.boxplot(x='institution_type', y='graduation_rate', data=df)
plt.title('Graduation Rate by Public vs. Private')
plt.xlabel('Institution Type')
plt.ylabel('Graduation Rate')
plt.show()
```



```
[290]: # Visualization: Student-to-Faculty Ratio by Institution Type
plt.figure(figsize=(8, 6))
sns.boxplot(x='institution_type', y='student_faculty_ratio', data=df)
plt.title('Student-to-Faculty Ratio by Institution Type')
plt.xlabel('Institution Type')
plt.ylabel('Student-to-Faculty Ratio')
plt.show()
```

<Figure size 800x600 with 0 Axes>



```
[212]: # Calculate the correlation between 'graduation_rate' and variables in the
↳ dataset

# Calculate the correlation matrix for the dataframe
correlation_matrix = df.corr()

# Extract the 'graduation_rate' column to see correlations with other variables
graduation_rate_correlation = correlation_matrix['graduation_rate'].
↳ sort_values(ascending=False)

# Print the correlations
print(graduation_rate_correlation)
```

```
graduation_rate          1.000000
out_of_state_tuition     0.596518
in_state_tuition         0.580855
total_cost_in_state      0.574065
percent_from_top_10      0.539111
percent_from_top_25      0.498203
board_costs              0.390641
```

```

percent_faculty_with_phd    0.287155
room_costs                  0.243991
applications_received       0.137240
applications_accepted       0.063971
additional_fees             0.050750
estimated_book_costs        0.005696
new_students_enrolled      -0.059394
full_time_undergrads       -0.096151
total_enrollment           -0.164624
estimated_personal_expenses -0.197224
acceptance_rate            -0.236663
student_faculty_ratio      -0.284587
part_time_undergrads       -0.297154
yield_rate                  -0.318371
Name: graduation_rate, dtype: float64

```

```

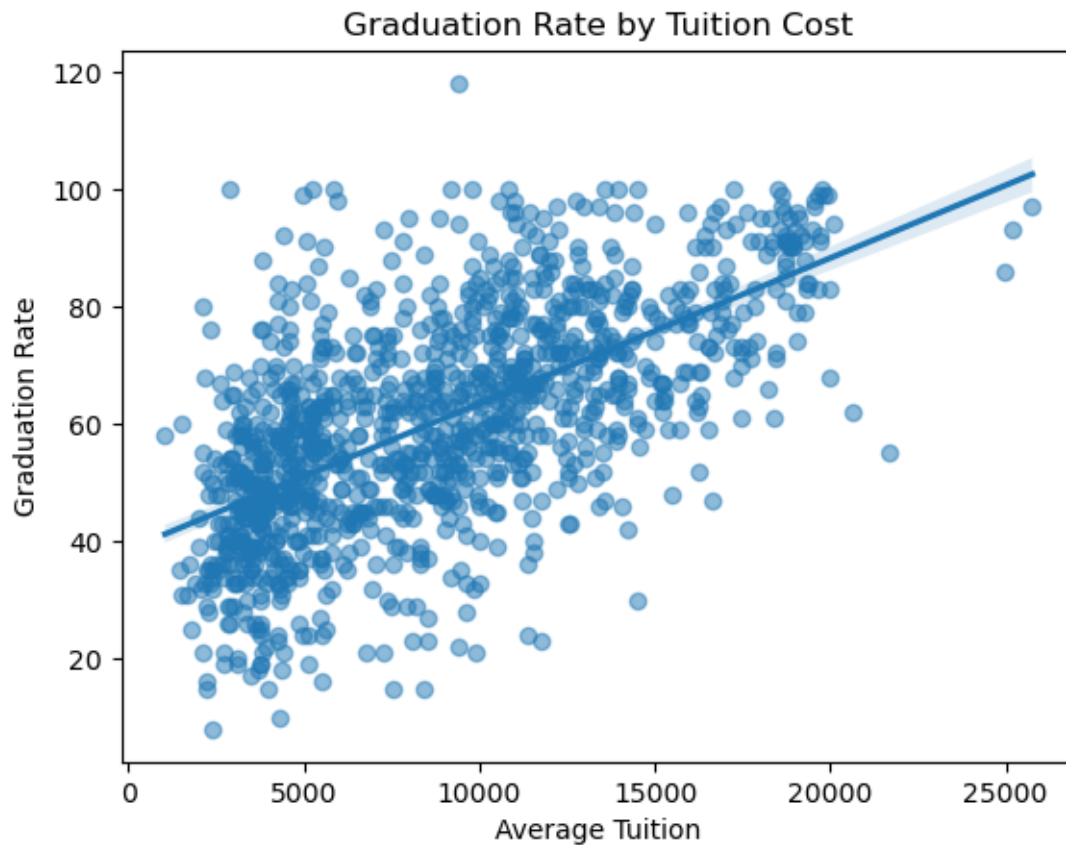
C:\Users\James\AppData\Local\Temp\ipykernel_29048\2859158846.py:2:
FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only valid
columns or specify the value of numeric_only to silence this warning.
    correlation_matrix = df.corr()

```


```

[ ]: sns.regplot(x='average_tuition', y='Graduation rate', data=df,
    ↪scatter_kws={'alpha':0.5})
plt.title('Graduation Rate by Tuition Cost')
plt.xlabel('Average Tuition')
plt.ylabel('Graduation Rate')
plt.show()

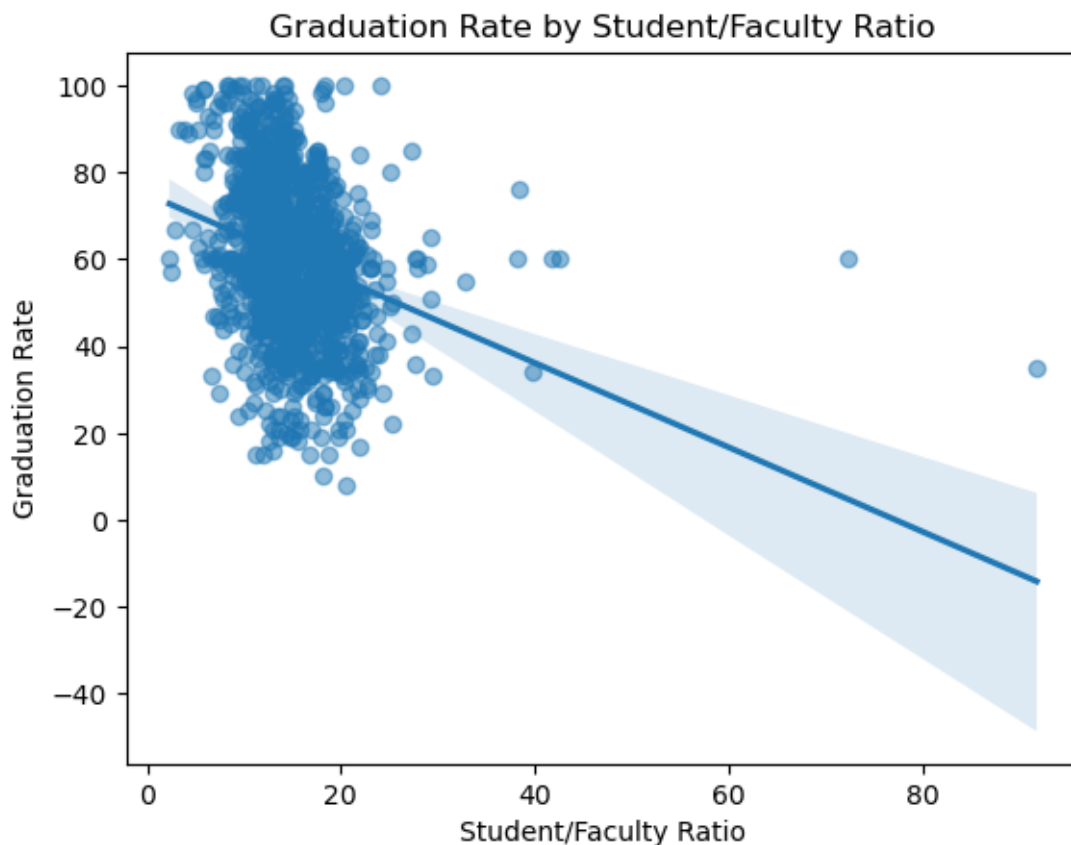
```



```
[ ]: sns.regplot(x='% fac. w/PHD', y='Graduation rate', data=df,
    ↳scatter_kws={'alpha':0.5})
plt.title('Graduation Rate by % Faculty w/PHD')
plt.xlabel('% Faculty w/PHD')
plt.ylabel('Graduation Rate')
plt.show()
```



```
[292]: sns.regplot(x='student_faculty_ratio', y='graduation_rate', data=df,
    ↪scatter_kws={'alpha':0.5})
plt.title('Graduation Rate by Student/Faculty Ratio')
plt.xlabel('Student/Faculty Ratio')
plt.ylabel('Graduation Rate')
plt.show()
```



```
[295]: # Generate Summary Statistics
df.describe()
```

```
[295]:
```

	applications_received	applications_accepted	new_students_enrolled	\
count	1171.000000	1171.000000	1171.000000	
mean	2661.401366	1838.443211	769.619129	
std	3409.315991	2188.810559	882.600645	
min	52.000000	36.000000	18.000000	
25%	687.000000	553.500000	233.000000	
50%	1450.000000	1086.000000	444.000000	
75%	3280.500000	2303.000000	945.000000	
max	48094.000000	26330.000000	7425.000000	

	percent_from_top_10	percent_from_top_25	full_time_undergrads	\
count	1171.000000	1171.000000	1171.000000	
mean	20.906917	44.493595	3647.936806	
std	18.119802	26.139218	4496.513985	
min	0.000000	0.000000	88.000000	
25%	9.000000	29.000000	967.500000	
50%	17.000000	46.000000	1803.000000	

75%	29.000000	63.000000	4475.000000
max	96.000000	100.000000	31643.000000

	part_time_undergrads	in_state_tuition	out_of_state_tuition \
count	1171.000000	1171.000000	1171.000000
mean	1047.527754	7909.197267	9301.883860
std	1648.558020	5277.480869	4098.244588
min	1.000000	480.000000	1044.000000
25%	128.000000	2615.500000	6161.000000
50%	446.000000	8190.000000	8734.000000
75%	1283.000000	11535.000000	11658.000000
max	21836.000000	25750.000000	25750.000000

	room_costs ...	percent_faculty_with_phd	student_faculty_ratio \
count	904.000000 ...	1171.000000	1171.000000
mean	2519.537611 ...	68.756618	14.881981
std	1151.862241 ...	17.536327	5.252689
min	500.000000 ...	8.000000	2.300000
25%	1714.250000 ...	57.000000	11.800000
50%	2200.000000 ...	71.000000	14.300000
75%	3010.000000 ...	82.000000	17.500000
max	7400.000000 ...	100.000000	91.800000

	graduation_rate	acceptance_rate	yield_rate	total_cost_in_state \
count	1171.000000	1171.000000	1171.000000	1171.000000
mean	60.513237	76.088792	45.041695	13391.611443
std	17.909663	14.911831	16.425049	6136.324488
min	8.000000	15.448631	9.975397	2521.000000
25%	49.000000	68.689057	33.477811	8411.000000
50%	60.000000	78.516903	42.250000	12906.000000
75%	73.000000	86.136346	53.475960	17875.500000
max	100.000000	100.000000	244.243421	29355.000000

	total_enrollment	total_cost	total_students	percent_full_time
count	1171.000000	1171.000000	1171.000000	1171.000000
mean	4695.464560	13391.611443	4695.464560	78.954166
std	5602.006605	6136.324488	5602.006605	16.390979
min	100.000000	2521.000000	100.000000	11.431412
25%	1229.500000	8411.000000	1229.500000	68.666242
50%	2323.000000	12906.000000	2323.000000	83.503836
75%	5730.000000	17875.500000	5730.000000	91.899260
max	38338.000000	29355.000000	38338.000000	99.941349

[8 rows x 24 columns]

1.6 Discussion:

Summary Statistics for the datasets were as follows (see table above): - Average applications received: 2661.4 - Average acceptance rate: 76% - Average total enrollment (full and part-time): 4,695.4 - Average total annual cost of attendance: \$13391.61

Analysis of the data revealed the following insights: - The Greatest number of universities were located in New York, Pennsylvania, and Texas. - The Average Total Annual Cost of Private Universities was OVER DOUBLE that of Public Universities, despite Public and Private Institutions being roughly just as selective. Despite this Graduation Rates at Private Universities are higher than Public (though this finding was not confirmed to be statistically significant using a hypothesis test, a task for further study). - Roughly 80% of the average student body are full-time students - Factors that most predicted high graduation rates were: - Tuition cost - Percent from Top 10%, 25% of their High School Class - Low Student-Faculty Ratios - Regarding graduation rates, percent of Faculty with PhD only showed a moderate effect on graduation rates, suggesting a potential cost-saving measure.

1.7 Conclusion:

This project demonstrates the power of data analytics in uncovering meaningful insights from complex datasets. By applying a range of data analysis techniques, it was possible to provide the consulting firm with a comprehensive understanding of the higher education landscape, enabling informed strategic advice for their clients.

Future directions for this work could involve deeper analysis with larger datasets, the integration of external data sources for richer insights, and the application of machine learning models to predict trends and outcomes in the education sector.